

# Representing Numbers of a Computer, Part 2

How computers store real numbers  
and the problems that result



1

## IEEE – Bitwise Storage Size



- The number of bits for the mantissa and exponent for the normal floating-point types are defined as:

Type	Sign, s	Exponent, c	Mantissa, f	Representation
Single 32-bit	1bit	8bits	23+1 bits	$(-1)^s \times 1.f \times 2^{c-127}$ Decimal: ~8s.f. $\times 10^{\pm 38}$
Double 64-bit	1bit	11bits	52+1 bits	$(-1)^s \times 1.f \times 2^{c-1023}$ Decimal: ~16s.f. $\times 10^{\pm 308}$

- there are also “Extended” versions of both the single and double types, allowing even more bits to be used. “Half Precision” (16-bit) also available for graphics, etc.
- the Extended types are not supported uniformly over a wide range of platforms; Single and Double are.



2




2


## IEEE Floating-point Discretisation

- This still cannot represent all numbers.
- E.g. with 5 bits for mantissa (including hidden bit):

- and in two dimensions you get something like:



epcc



3

## 32-bit and 64-bit floating point

- Conventionally called single and double precision
  - C, C++ and Java: `float` (32-bit), `double` (64-bit)
  - Fortran: `real` (32-bit), `double precision` (64-bit)
    - or `real(kind(1.0e0))`, `real(kind(1.0d0))`
    - or `real(kind=4)`, `real(kind=8)`
  - Nothing to do with 32-bit / 64-bit operating systems!!!
- Single precision accurate to ~8 significant figures
  - E.g. 3.2037743E+03
- Double precision accurate to ~16 significant figures
  - E.g. 3.203774283170437E+03
- Fortran usually knows this when printing default format
  - C and Java often don't
  - depends on compiler

epcc

4



4

## Limitations

- Numbers cannot be stored exactly
  - gives problems when they have very different magnitudes
- E.g. 1.0E-6 and 1.0E+6
  - no problem storing each number separately, but when adding:

$$0.000001 + 1000000.0 = 1000000.000001 = 1.000000000001E6$$

- in 32-bit will be rounded to 1.0E6

- So

$$(0.000001 + 1000000.0) - 1000000.0 = 0.0 \quad \times$$

$$0.000001 + (1000000.0 - 1000000.0) = 0.000001 \quad \checkmark$$

- FP arithmetic is commutative:  $A + B = B + A$
- ...but not in general associative  $(A + B) + C \neq A + (B + C)$



5



5

## Example I

```

program recurrence
  implicit none
  real (kind=4) :: s23
  real (kind=8) :: d23
  real (kind=16) :: q23
  integer :: i

  s23 = 2.0 / 3.0
  d23 = 2.0_8 / 3.0_8
  q23 = 2.0_16 / 3.0_16

  do i = 1, 18
    s23 = s23 / 10 + 1
    d23 = d23 / 10 + 1
    q23 = q23 / 10 + 1
    write(*,*) s23, d23, q23
  end do

  do i = 1, 18
    s23 = (s23 - 1) * 10
    d23 = (d23 - 1) * 10
    q23 = (q23 - 1) * 10
    write(*,*) s23, d23, q23
  end do
end program recurrence
  
```

- start with  $\frac{2}{3}$
- Set up single, double, quadruple
- divide by 10 add 1
- repeat many times (18)
- subtract 1 multiply by 10
- repeat many times (18)

6



6

[illegible]

7

7

## The result: Two thirds

Single precision  
fifty three billion!

## Double precision fifty!

Quadruple precision **loses** information about two-thirds after 18<sup>th</sup> decimal place



8

8

## Example II – order matters!

```
#include <iostream>

template <typename T>
void order(const char* name) {
    T a, b, c, x, y;

    a = -1.0e10;
    b = 1.0e10;
    c = 1.0;
    x = (a + b) + c;
    y = a + (b + c);

    std::cout << name << ": x = " << x << ", y = " << y << std::endl;
}

int main()
{
    order<float>(" float");
    order<double>("double");
    return 0;
}
```

This code adds three numbers together in a different order. Single and double precision.

$$x = (-1.0 \times 10^{10} + 1.0 \times 10^{10}) + 1.0$$
$$y = -1.0 \times 10^{10} + (1.0 \times 10^{10} + 1.0)$$

What is the answer?



9



9

## The result. One

```
$ clang++ -O0 order.cpp -o order
$ ./order
float: x = 1, y = 0 ✗
double: x = 1, y = 1 ✓
```



10



10

## Example III: Gauss

- C. 1785 AD in what is now Lower Saxony, Germany
  - School teacher sets class a problem
  - Sum numbers 1 to 100
  - Nine year old boy quickly has the answer

$$S_n = \sum_{i=1}^n i = \frac{n}{2}(n+1)$$

$$S_{100} = \frac{100}{2}(100+1) = 5050$$



Carl Friedrich Gauss  
(C.1840 AD)

## Summing numbers

```
#include <stdio.h>

int main() {
    int i, m;
    float sum_up, sum_down;
    int n = 100;

    for (m = 0; m < 3; ++m) {
        sum_up = 0;
        for (i = 1; i <= n; ++i) {
            sum_up += i;
        }

        sum_down = 0;
        for (i = n; i >= 1; --i) {
            sum_down += i;
        }

        printf("Gaussian sum up to %5d: %11.1f %11.1f %9.1f\n",
            n, sum_up, sum_down, n*(n+1)/2);
        n *= 10;
    }
}
```

sums numbers to 100, 1000, 10000  
performs sum low-to-high and high-  
to-low in single precision

## The result: Gauss' sum

```
$ clang gauss.c -o gauss
$ ./gauss
Gaussian sum up to 100:      5050.0      5050.0      5050
Gaussian sum up to 1000:    500500.0    500500.0    500500
Gaussian sum up to 10000:  50002896.0  50009072.0  50005000
```

In single precision summing numbers 1 to 10000 produces the **wrong** answer  
high-to-low and low-to-high produce **different** wrong answers

What happens when in parallel  
same calculation, different numbers of processors!



13



13

## Special Values

- In floating point numbers, zero is treated specially
  - corresponds to all bits being zero (except the sign bit)
    - Zero can have a sign
    - Both +0.0 and -0.0 equate to be the same in calculations
    - $1.0/(-0.0) = -\infty$  and  $1.0/(+0.0) = +\infty$
- There are other special numbers
  - infinity: which is usually printed as "Inf"
  - Not a Number: which is usually printed as "NaN"
- These also have special bit patterns



14



14

## Infinity and Not a Number

- Infinity is usually generated by dividing any finite number by 0.
  - although can also be due to numbers being too large to store
  - some operations using infinity are well defined, e.g.  $-3/\infty = -0$
- NaN is generated under a number of conditions:
  - $\infty + (-\infty)$ ,  $0 \times \infty$ ,  $0/0$ ,  $\infty/\infty$ ,  $\sqrt{x}$  where  $x < 0.0$
  - most common is the last one, e.g.  $x = \text{sqrt}(-1.0)$
- Any computation involving NaNs returns NaN.
  - there is actually a whole set of NaN binary patterns, which can be used to indicate why the NaN occurred.



15



15

## IEEE Special Values

Exponent, e (unshifted)	Mantissa, f	Represents
000000...	0	$\pm 0$
000000...	$\neq 0$	$0.f \times 2^{(1-bias)}$ [subnormal]
$000... < e < 111...$	Any	$1.f \times 2^{(e-bias)}$
111111...	0	$\pm \infty$
111111...	$\neq 0$	NaN

- Most numbers are in standard form (middle row)
  - have already covered zero, infinity and NaN
  - but what are these “subnormal numbers” ???



16



16



## Range of Single Precision

- Have 8 bits for exponent, 1+23 bits for mantissa
  - unshifted exponent can range from 0 to 255 (bias is 127)
  - smallest and largest values are reserved for zero, subnormal (see later) and infinity or NaN
  - unshifted range is then 1 to 254, shifted is -126 to 127

- Largest number:

$$1.111111111111111111111111 \times 2^{127} \\ \sim 2 \times 2^{127} = 2^{128} \sim \mathbf{3.4 \times 10^{38}}$$

- Smallest number

$$= 2^{-126} \sim 1.2 \times 10^{-38}$$

- But what is the case of the zero exponent (non-zero mantissa) reserved for ...?



17



17

# IEEE Subnormal Numbers

- Standard IEEE has mantissa normalised to  $1.xxx$
- But, normalised numbers can give  $x - y = 0$  when  $x \neq y$ !
  - consider  $1.10 \times 2^{-E_{min}}$  and  $1.00 \times 2^{-E_{min}}$  where  $E_{min}$  is smallest exponent
  - upon subtraction, we are left with  $0.10 \times 2^{-E_{min}}$ .
  - in normalised form we get  $1.00 \times 2^{-E_{min}-1}$ :
    - this cannot be stored because the exponent is too small.
    - when normalised it must be flushed to zero, giving a gap in this region.
  - thus, we have  $x \neq y$  while at the same time  $x - y = 0$ !
- Thus, the smallest exponent is set aside for *subnormal* (or *denormal*) numbers, beginning with  $0.f$  (not  $1.f$ ).
  - can store numbers smaller than the normal minimum value
    - but with reduced precision in the mantissa
  - ensures that  $x = y$  when  $x - y = 0$  (also called *gradual underflow*)



18



18

## Subnormal Example

- Consider the single precision bit patterns:
  - mantissa: 0000100....
  - exponent: 00000000
- Exponent is zero but mantissa is non-zero
  - a subnormal number
  - value is  $0.0000100... \times 2^{-126} \sim 2^{-5} \times 2^{-126} = 2^{-131} \sim 3.7\text{E-}40$
- Smaller than normal minimum value
  - but we lose precision due to all the leading zeroes
  - smallest possible number is  $2^{-23} \times 2^{-126} = 2^{-149} \sim 1.4\text{E-}45$



19



19

## Exceptions

- May want to terminate calculation if any special values occur
  - could indicate an error in your code
- Can usually be controlled by your compiler
  - default behaviour can vary
  - E.g. some systems terminate on NaN, some continue
- Usual action is to terminate and dump the core



20



20

## IEEE Arithmetic Exceptions

Exception	Result
Overflow	$\pm\infty$ , $f = 11111\dots$
Underflow	$0, \pm 2^{-bias}$ , [subnormal]
Divide by zero	$\pm\infty$
Invalid	NaN
Inexact	$round(x)$

- It is not necessary to catch all of these.
  - inexact occurs extremely frequently and is usually ignored
  - underflow is also usually ignored
  - you probably want to catch the others

## IEEE Rounding

- We wish to add, subtract, multiply and divide.
  - E.g. Addition of two decimal numbers, given to 4 s.f.:

$$\begin{array}{rclcl}
 0.1241 \times 10^{-1} & + & 0.2815 \times 10^{-2} & = & \\
 0.1241 \times 10^{-1} & + & 0.02815 \times 10^{-1} & = & 0.15225 \times 10^{-1} \\
 \hline
 \text{But can only store 4 significant figures:} & & & & 0.1522 \times 10^{-1} \\
 & & & & \text{or} \\
 & & & & 0.1523 \times 10^{-1}
 \end{array}$$

- In essence:
  - we shift the decimal point on one input as required,
  - perform fixed point arithmetic,
  - renormalise the number by shifting the decimal point again.
- But what do we do with that 5?
  - do we round up, round down, truncate, ...

## IEEE Rounding Modes

- We can choose from several rounding types:
  - there are four types of rounding for arithmetic operations.
    - Round to nearest: e.g. -0.001298 becomes -0.00130.
    - Round to zero: e.g. -0.001298 becomes -0.00129.
    - Round to +infinity: e.g. -0.001298 becomes -0.00129.
    - Round to -infinity: e.g. -0.001298 becomes -0.00130.
  - but how can we ensure the rounding is done correctly?
- **Guard digits:**
  - calculations are performed at slightly greater precision on the CPU, and then stored in standard IEEE floating-point numbers.
  - usually uses three extra binary bits to ensure correctness.
  - Guarantees an FP operation gives correct result up to rounding
    - answer will still be inaccurate due to rounding and rounding errors will accumulate
- **Your compiler may be able to change the mode**
  - Round to nearest is default



23



23

## Implementations: C & Fortran

- Most C and Fortran compilers are fully IEEE 754 compliant.
  - compiler switches are used to switch on exception handlers.
  - these may be very expensive if dealt with in software.
  - you may wish to switch them on for testing (except **inexact**), and switch them off for production runs.
- But there are more subtle differences.
  - Fortran always preserves the order of calculations:
    - $A + B + C = (A + B) + C$ , always.
  - C compilers are free to modify the order during optimisation.
    - $A + B + C$  may become  $(A + B) + C$  or  $A + (B + C)$ .
    - Usually, switching off optimisations retains the order of operations.
- **Complex numbers usually stored as pair of floating point numbers**
  - C (ISO C99) and Fortran support for going between real and complex



24



24

## Implementations: Java

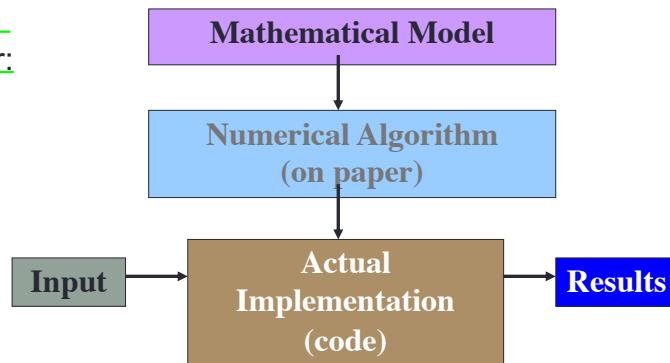
- In summary:
  - Java only supports round-to-nearest.
  - Java does not allow users to catch floating-point exceptions.
  - Java only has one NaN.
- All of this is technically a bad thing
  - these tools can be used to test for instabilities in algorithms
  - this is why hardcore numerical scientists don't like Java very much
  - however, Java also has some advantages over, say, C
    - forces explicit casting
    - you can use the strictfp modifier to ensure that the same bytecode produces identical results across all platforms.

## Other Precisions

- Half-precision
  - Uses 16-bit
  - Used where high precision is not needed but storage is at a premium
  - Encountered in graphics – e.g. OpenGL, GPUs
- Quadruple (or quad) precision
  - mantissa 112+1 bits, exponent 15 bits
  - ~32 significant figures
- Higher precision?
  - “Multiple precision” libraries exist but SLOW!
  - Some computer algebra packages (e.g. Maple, Mathematica) allow arbitrary precision but these tend to be slow at computation anyway

## More on errors...

- We want fast, accurate and stable algorithms.
  - But all numerical algorithms produce inaccurate results.
- There are three sources of error:
  - Truncation
  - Rounding
  - Data



- And there is always a trade-off between speed & accuracy.



27



27

## Data Errors

- Poor data.
  - It is very difficult to write code that will **compensate** for errors in your input data.
- Badly stored data:
  - The simplest failure can occur on the loading and saving of data.



Inputted and stored as integer or floating-point.

Output from internal representation to decimal (usually).

- Input too fine?
- Output too coarse?



28



28

## Rounding errors: Diagnosis

- How can you tell if your algorithm is suffering due to rounding errors?
  - There is no **sure-fire** way, but...
  - If you make small changes to the input data, do you see:
    - Wild (chaotic?) variations in the output?
    - No change in the output at all?
  - Do you get the same answer if you change the level of precision?  
E.g. 32-bit to 64-bit.
  - Are you failing to catch important FP exceptions?
  - Does changing the rounding-mode give you very different answers?



29



29

## Solutions

- Find a better algorithm!
  - Calculating the area of a triangle – e.g. recent practical!
  - However, this could be impossible in a research situation.
- Use a greater level of precision.
  - Up to 128-bit numbers supported on most platforms (slow).
  - Arbitrary precision calculations may also be an option.
  - Beware! The rounding error is not always proportional to the numerical precision!
- Identify the **range** of inputs for which your algorithm is reliable.
  - And stick to it!



30



30

## Catastrophic cancellation

- Even subtraction can be dangerous...
  - When the two numbers are nearly equal:
  - E.g.  $10000000.23 - 10000000.22 = 0.01$
- If both are known precisely, then all is well.
  - We can be sure that the value of 0.01 is free from error.
- If the difference between the two numbers arises from rounding errors, the relative error is magnified.
  - E.g. perhaps the above numbers were supposed to be equal and the small difference was due to rounding error.
  - Result should then have been zero
  - Original values were correct to 9 s.f. (tiny relative error)
  - The relative error in the result is now infinite!
    - Relative error =  $(0.0 - 0.01)/0.0$

## Truncation Errors

- Rounding errors can be overcome by clever design and using sufficiently large resources.
- This is not true of truncation errors:
  - Due to the differences between the mathematical model and the numerical algorithm.
    - Independent of implementation or precision.
    - You are solving a different system to the one you want.
  - Examples:
    - Truncated Taylor-series.
      - The sin function may be calculated as:
        - $\sin(x) = x - \frac{x^3}{3!} + \frac{x^5}{5!} - \frac{x^7}{7!} + \mathcal{O}(x^9)$
      - How does this approximation affect the accuracy of the algorithm?
    - Discretisation of the problem domain...



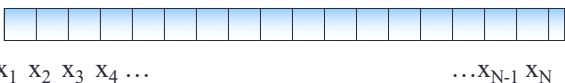
## Discretisation Example

- Heat diffusion over a 1D metal bar:

- Ends held at two different temperatures,  $T_1$  and  $T_2$ .

$T_1$    $T_2$

- This must be discretized to solve it numerically:

$T_1$    $T_2$   
 $x_1 \ x_2 \ x_3 \ x_4 \ \dots \ \dots x_{N-1} \ x_N$

- You are solving the diffusion equation for a discrete system, not the system you originally wanted. Is this OK?
- The difference between the results for the original system and the numerical one is the truncation error.

## Summary

- Floating point numbers defined in IEEE 754 standard
  - defines storage format
  - can be single (32-bit) and double (64-bit) precision
  - and the result of all arithmetical operations
- All real calculations suffer from rounding errors
  - important to choose an algorithm where these are minimised
- Practical exercise illustrates the key points