

Numerical Algorithms for HPC

Linear Algebra Libraries:
LAPACK and ScaLAPACK



1

Linear algebra libraries

- Linear Algebra
- Matrix types
- Serial Libraries
 - BLAS
 - LINPACK
 - LAPACK
 - LU Factorisation
- Parallel libraries
 - ScaLAPACK



2

2

Linear algebra libraries

- Linear Algebra is a well constrained problem
 - can define a small set of common operations
 - implement them robustly and efficiently in a library
 - mainly designed to be called from Fortran (see later ...)
- Often seen as the most important HPC library
 - e.g. LINPACK benchmark
- Linear algebra is unusually efficient
 - LU decomposition has $O(N^3)$ operations for $O(N^2)$ memory loads



3



3

Matrix types

- Matrices generally classified as either *sparse* or *dense*
 - We will deal with sparse matrices later
- Rectangular matrices
 - correspond to different number of equations from unknowns
 - system can be either under- or over-determined
- Matrices may have symmetry about the diagonal:
 - Symmetric (real matrices): $a_{ij} = a_{ji} \rightarrow A^T = A$
 - Hermitian (complex matrices): $a_{ij}^* = a_{ji} \rightarrow A^{T*} = A^\dagger = A$

$$\begin{pmatrix} 1 & 2 \\ 2 & 3 \end{pmatrix}$$

Symmetric

$$\begin{pmatrix} 1 & 2+i \\ 2-i & 3 \end{pmatrix}$$

Hermitian




4




4

Matrix structures


- Many matrices have a regular structure



Upper
triangular



Band
diagonal



Block
diagonal

- Can exploit regular structure or symmetry for efficiency
 - eg special form of LU decomposition for tridiagonal matrices



5




5

BLAS

- Basic Linear Algebra Subprograms
 - Level 1: vector-vector operations (e.g. $x \cdot y$)
 - Level 2: matrix-vector operations (e.g. Ax)
 - Level 3: matrix-matrix operations (e.g. AB)
 (x, y vectors, A, B matrices)

- Example: SAXPY routine



y is replaced “in-place” with $a x + y$



6



6

SAXPY

- In Fortran

```
call SAXPY(n,a,x,incx,y,incy)
...
ix = 1; iy = 1
do i = 1, n
    y(iy) = y(iy) + a * x(ix)
    ix = ix + incx; iy = iy + incy
end do
```



7



7

LINPACK library



- <http://www.netlib.org/linpack/>
- LINPACK "Top 500" benchmark: HPL (High Performance LINPACK) is standard HPC performance metric
- Possible to achieve performance close to theoretical peak
- LINPACK *benchmark* means LU factorisation
- Collection of Fortran subroutines that analyse and solve linear equations and linear least-squares problems.
- Solves linear systems whose matrices are general, banded, symmetric indefinite, symmetric positive definite, triangular and tridiagonal square.
- Package can also compute QR and singular value decompositions (SVD) of rectangular matrices and applies them to least-squares problems.




8



8

Why is performance so high? LU factorisation

- Look at step in LU decomposition
 - consider an in-place decomposition
 - for each column k
 - scale the sub-diagonal column below a_{kk} by $1/a_{kk}$
 - for all columns j to the right, subtract a_{kj} times above
 - e.g. for $k = 1$



- These are vector operations over columns
 - scaling a vector x by a (a scalar variable): $x_i \rightarrow ax_i$
 - updating a second vector y by the above: $y_i \rightarrow ax_i + y_i$



9



9

LINPACK implementation

- Based around level 1 BLAS
 - e.g. look at LU decomposition

```

t = -1.0e0/a(k,k)
call SSCAL(n-k,t,a(k+1,k),1)
...
do j = k+1,n
  call SAXPY(n-k,-a(k,j),a(k+1,k),1,a(k+1,j),1)
end do

```

- Very efficient on early HPC vector machines
 - not nearly so efficient on modern cache-based systems



10



10

LAPACK

- LAPACK is built on top of BLAS libraries
 - Most of the computation is done with the BLAS libraries
- Original goal of LAPACK was to run efficiently on shared memory and multi-layered systems
 - Spend less time moving data around!
- LAPACK attempts to use BLAS 3 instead of BLAS 1
 - matrix-matrix operations more efficient than vector-vector
- Illustrates trend to layered numerical libraries
 - allows for portable performance libraries
 - efficient implementation of BLAS 3 leads immediately to efficient implementation of LAPACK
 - porting LAPACK becomes a straightforward exercise



11



11

BLAS/LAPACK naming conventions

- Routines generally have a name of up to 6 letters, e.g. DGESV,
- Initial letter
 - S: Real C: Complex
 - D: Double Precision Z: Double Complex or COMPLEX*16
- For level 2 and 3 routines, 2nd and 3rd letter refers to matrix type
 - GE: matrices are general rectangular
 - i.e. could be unsymmetric, not necessarily square
 - HE: (complex) Hermitian
 - SY: symmetric
 - TR: triangular
 - BD: bidiagonal
 - etc. ~30 in total
- E.g. SGESV: Single precision, general matrix solver (solves $Ax = b$)



12



12

LU factorisation

- LU factorisation
 - `call SGETRF(M, N, A, LDA, IPIV, INFO)`
 - does an in-place LU factorisation of M by N matrix A
 - we will always consider the case $M = N$
 - A can actually be declared as `REAL A(NMAX,MMAX)`
 - routine operates on $M \times N$ submatrix
 - must tell the library the Leading Dimension of A , i.e. set `LDA=NMAX`
 - `INTEGER IPIV(N)` returns row permutation due to pivoting
 - error information returned in the integer `INFO`



13



13

Solving: Forward/backward substitution

- Forward / backward substitution
 - `call SGETRS(TRANS,N,NRHS,A,LDA,IPIV,B,LDB,INFO)`
 - expects a factored A and `IPIV` from previous call to `SGETRF`
 - solves for multiple right-hand-sides, i.e. B is $N \times NRHS$
 - we will only consider `NRHS=1`, i.e. RHS is the usual vector \underline{b}
 - solution x is returned in \underline{b} (i.e. original \underline{b} is destroyed)
- Options exist for precise form of equations
 - specified by character variable `TRANS`
 - 'N' (Normal), 'T' (Transpose) or 'C' (hermitian Conjugate)

$$\begin{array}{ccc}
 \nwarrow & \nearrow & \nearrow \\
 \underline{A} \underline{x} = \underline{b} & \underline{A}^T \underline{x} = \underline{b} & \underline{A}^\dagger \underline{x} = \underline{b}
 \end{array}$$



14



14

Calling Fortran libraries from C

- A number of issues
 - storage order of matrices
 - calling by reference / calling by value
 - character variables
 - subroutine names

System Dependent
- C arrays are transposed w.r.t. Fortran
 - could choose to store all matrices in transpose format
 - but may simply be able to specify **TRANS= 'T'** where appropriate
- Fortran *a/ways* expects pass-by-reference
 - must assign C constants to variables, eg **one = 1;**
 - pass the pointer **&one** to the subroutine



15



15

Calling LAPACK from C

- Fortran


```
call SGETRS (TRANS, N, NRHS, A, LDA, IPIV, B, LDB, INFO)
```
- Easiest to write a wrapper for C, e.g:


```
int sgetrs(char trans, int n, int nrhs,
float *a, int lda, int *ipiv, float *b, int ldb)
{
    int info;
    sgetrs_(&trans, &n, &nrhs, a, &lda, ipiv, b, &ldb, &info);
    return(info);
}
...
info = sgetrs('t', n, 1, &(a[0][0]), NMAX, ipiv, x, NMAX);
```

C requires the following libs when linking:
`-llapack -lblas -lpqftnrtl -lrt`



16



16

Sparse matrices

- Direct methods easiest for structured matrices
 - E.g. tridiagonal, pentadiagonal, block-diagonal
 - support in LAPACK for some well-structured sparse cases
- General sparse matrices
 - difficult to code efficiently due to “fill-in”
 - LU factors will have non-zero entries even where A was zero
 - some specialist libraries, e.g. Harwell Sparse Matrix Library
- In general, iterative methods are used
 - see later



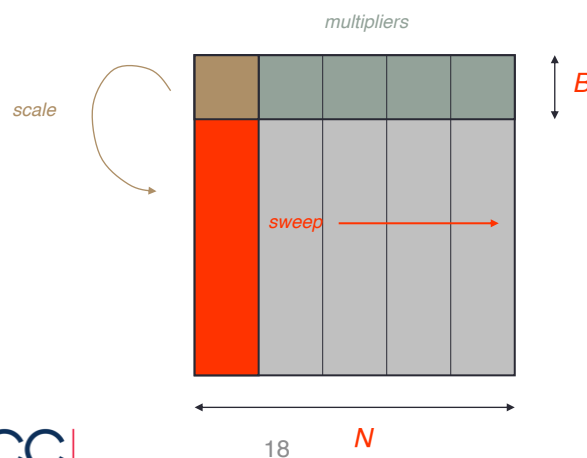
17



17

Blocked linear algebra

- LAPACK achieves performance by *blocking*
- Operate on $B \times B$ sub-matrices and not scalars




18



18

Basic operations



LU factorisation and solution

BLAS 3:
SGEMM

epcc


19



19

Parallelisation

- Clear opportunities for parallelism
 - multiple independent **saxpy** or **SGEMM** operations
 - major problem is load balance, on four processors



- No simple block distribution is appropriate
 - But clear we must use block-cyclic in at least one dimension

epcc

20



20

ScaLAPACK introduction

- ScaLAPACK allows us to run LAPACK-like routines *in parallel*
- Routines written to resemble equivalent LAPACK routines
 - e.g. dgesv → pdgesv
- Assumes matrices are laid out in 2D block-cyclic form
 - In contrast to sparse matrices - usually 1D decomposition
- Built on top of
 - LAPACK (Linear algebra library)
 - PBLAS (distributed memory version of Level 1, 2 and 3 BLAS)
 - BLACS (Basic Linear Algebra Communication Subprograms)
 - MPI
- As with LAPACK, written in Fortran 77 but interfaces to Fortran 90/C/C++




21



21

Hierarchy



<http://www.netlib.org/scalapack/>



22

ScaLAPACK

- ScaLAPACK follows similar format to MPI
 - BLACS “context” being the equivalent of an MPI communicator
 - Need several set-up and finalise routines
- Matrices/vectors completely distributed over processors and described using an *array descriptor*
- Set up a 2D processor grid
 - Routines provided to determine processor position in grid and local matrix size
- Data distributed in a “block cyclic” fashion with block size (referred to as “blocking factor”)
 - ScaLAPACK describes this as “complicated but efficient”



23



23

2D Block cyclic

- Situation can be simple – e.g. 8×8 matrix with 2×2 processor grid...
- ...or by making processor blocks smaller can be more complicated...

P ₀	P ₀	P ₀	P ₀	P ₁	P ₁	P ₁	P ₁
P ₀	P ₀	P ₀	P ₀	P ₁	P ₁	P ₁	P ₁
P ₀	P ₀	P ₀	P ₀	P ₁	P ₁	P ₁	P ₁
P ₀	P ₀	P ₀	P ₀	P ₁	P ₁	P ₁	P ₁
P ₂	P ₂	P ₂	P ₂	P ₃	P ₃	P ₃	P ₃
P ₂	P ₂	P ₂	P ₂	P ₃	P ₃	P ₃	P ₃
P ₂	P ₂	P ₂	P ₂	P ₃	P ₃	P ₃	P ₃
P ₂	P ₂	P ₂	P ₂	P ₃	P ₃	P ₃	P ₃

P ₀	P ₀	P ₁	P ₁	P ₀	P ₀	P ₁	P ₁
P ₀	P ₀	P ₁	P ₁	P ₀	P ₀	P ₁	P ₁
P ₂	P ₂	P ₃	P ₃	P ₂	P ₂	P ₃	P ₃
P ₂	P ₂	P ₃	P ₃	P ₂	P ₂	P ₃	P ₃
P ₀	P ₀	P ₁	P ₁	P ₀	P ₀	P ₁	P ₁
P ₀	P ₀	P ₁	P ₁	P ₀	P ₀	P ₁	P ₁
P ₂	P ₂	P ₃	P ₃	P ₂	P ₂	P ₃	P ₃
P ₂	P ₂	P ₃	P ₃	P ₂	P ₂	P ₃	P ₃



24



24

2D block cyclic distribution - example

- Global matrix size: 9×9
- Block size: 2×2
- No. processors: 6
- Processor grid: 2×3
- Local matrix sizes
 - P_0 : $5 \times 4 = 20$
 - P_1 : $5 \times 3 = 15$
 - P_2 : $5 \times 2 = 10$
 - P_3 : $4 \times 4 = 16$
 - P_4 : $4 \times 3 = 12$
 - P_5 : $4 \times 2 = 8$
- Routines exist to calculate local sizes!

P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁




25



25

2D block cyclic applied to LU factorisation



P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₃	P ₃	P ₄	P ₄	P ₅	P ₅	P ₃	P ₃	P ₄
P ₀	P ₀	P ₁	P ₁	P ₂	P ₂	P ₀	P ₀	P ₁

- This decomposition allows a reasonable load balance



26



26

ScaLAPACK: Example solving $Ax=b$

```

...
n=64; nrhs=1; nprow=2; npc=3; mb=nb=2;
...
CALL BLACS_GET(-1,0,ctxt)
CALL BLACS_GRIDINIT(ctxt, 'Row-major', nprow, npc)
CALL BLACS_GRIDINFO(ctxt, nprow, npc, myrow, mycol)
...
num_rows_local = NUMROC(n, nb, myrow, 0, nprow)
num_cols_local = NUMROC(n, nb, mycol, 0, npc)
...
Allocate(A(num_rows_local, num_cols_local), b(num_rows_local),
ipiv(num_rows_local+nb))
...
IF(MYROW.EQ.-1) Skip computation!
...
CALL DESCINIT(desca, n, n, mb, nb, rsrc, csrc, ctx, llda, info)
CALL DESCINIT(descb, n, nrhs, nb, nrhs, rsrc, csrc, ctx, lldb, info)
...
call PDGETRF(n, n, A, 1, 1, desca, ipiv, info)
call PDGETRS('N', n, 1, A, 1, 1, desca, ipiv, b, 1, 1, descb, info)
...
WRITE(*,*) ...Results.....
...
CALL BLACS_EXIT(0)

```



L05-LAPACK-ScaLAPACK

