

# Numerical Algorithms for HPC

Solution of Boundary Value Problems



1

## Overview of Lecture

- Overview of linear solving methods
- Relaxation methods
  - Jacobi algorithm
  - testing for convergence
  - Gauss Seidel
  - over-relaxation
- Notes
  - parallelisation
  - non-linear equations
- Pollution problem
  - solution using relaxation methods
  - 2D equations including wind



2



2

## Many methods for solving $Au = b$

- Direct methods
  - give the solution after a fixed number of operations
    - Gaussian elimination
    - LU factorisation
- Relaxation methods (this lecture)
  - gradually improve solution, starting from an initial guess
  - stop when the answer is sufficiently accurate
  - simple to implement but may be slow to converge on solution
    - or may fail completely!
- Krylov subspace methods (following lectures)
  - iterative (like relaxation methods) but more sophisticated
  - harder to implement but more efficient and reliable



3



3

## Why not use Direct Methods?

- Direct methods explicitly operate on the matrix  $A$ 
  - e.g. decompose it into  $L$  and  $U$  factors
- For PDEs,  $A$  is very sparse indeed
  - may contain 99% zeros so clearly we use compressed storage
  - we want to take advantage of this when we solve equations
- Difficult to exploit sparsity for direct methods
  - e.g.  $L$  and  $U$  may be dense even though  $A$  is sparse
  - for large systems of equations, we may run out of memory!
- Relaxation and Krylov methods (see later) exploit sparsity
  - relaxation methods operate on the equations not the matrix
  - Krylov methods comprise mostly matrix-vector multiplications
    - can write efficient routines to do  $y = Ax$  when  $A$  is sparse
  - start to show the solution earlier during process of solving



4



4

## Relaxation vs Matrix Methods

- Operate directly on the difference equations
  - can forget (almost!) all about the matrix representation  $Au = b$  for this lecture
  - it turns out that relaxation methods can usefully be understood in terms of matrix-vector operations (not immediately obvious)
    - See lecture on “Matrix Splitting Techniques” later
- For illustrative purposes, look at 1D problem
  - for simplicity with no wind
  - exercise will involve extending this to the 2D problem
    - quite straightforward in practice



5



5

## Relaxation Methods

- 1D diffusion equations are
$$-u_{i-1} + 2u_i - u_{i+1} = 0 \quad i = 1, 2, \dots, N$$
- Equivalently:  $u_i = 1/2 (u_{i-1} + u_{i+1})$ 
  - why not make an initial guess at the solution
  - then loop over each lattice point  $i$  and set  $u_i = 1/2 (u_{i-1} + u_{i+1})$
  - i.e. we solve the equation exactly at each point in turn
- Updating  $u_i$  spoils the solution we just got for  $u_{i-1}$ 
  - so simply iterate the whole process again and again ...
  - ... and hope we eventually get the right answer!
- This is called the Jacobi Algorithm
  - the simplest possible relaxation method



6



6

## Jacobi Algorithm

- Use superscript  $n$  to indicate iteration number
  - $n$  counts the number of times we update the *whole* solution
  - equivalent to computer time

- Jacobi algorithm for diffusion equation is:

$$u_i^{(n+1)} = 1/2 (u_{i-1}^{(n)} + u_{i+1}^{(n)})$$

- Each iteration, calculate  $u^{(n+1)}$  in terms of  $u^{(n)}$ 
  - don't need to keep copies of all the previous solutions
  - only need to remember two solutions at any time:  $u$  and  $u_{\text{new}}$ 
    - corresponding to iterations  $n$  and  $n + 1$



7



7

## Jacobi Pseudo-Code

```
declare arrays:    u(0, 1, ..., M+1)
                  unew(0, 1, ..., M+1)

initialise: set boundaries: u(0)    = fixed value  $u_0$ 
                  u(M+1) = fixed value  $u_{M+1}$ 
                  initial guess: u(1, 2, ..., M) = guess value

loop over n = 1, 2, ...

    update:      loop over internal points: i = 1, 2, ... M
                  unew(i) = 0.5*( u(i-1) + u(i+1) )
                  end loop over i

    copy back: u(1, 2, ..., M) = unew(1, 2, ..., M)

end loop over n
```



8







8

## Implementation Notes

- Array declarations
  - Fortran: `real, dimension(0:M+1) :: u`
  - Java: `float[] u = new float[M+2];`
  - C: `float u[M+2];`
- Arrays explicitly contain boundaries  $u_0$  and  $u_{M+1}$ 
  - we set them according to boundary conditions
    - but we NEVER update them!
  - e.g. when we copy  $u_{\text{new}}$  back to  $u$ , only copy internal values
  - in pseudo-code, boundary values for  $u_{\text{new}}$  are never set
    - complete solution is therefore only ever present in  $u$
    - might be more elegant to set boundaries in  $u_{\text{new}}$  as well
- What to choose for initial guess  $u_i^{(0)}$ ?
  - for a simple implementation just set interior values to zero

## Progress of Solution



## When to Stop the Iterative Loop

- The solution *appears* to be getting better
  - must quantify this!
- For dense systems we used the residual
  - we tried to solve  $Ax = b$ , so  $r = b - Ax$  should be a zero vector
  - in practice, there is a numerical error in solution of each equation
  - error in equation  $i$  is the value of  $r_i$ 
    - Norm of residual is computed from the sum of the squares of  $r_i$
  - Can calculate residue as before:  $\frac{\|r\|_2}{\|b\|_2}$
- Can do the same thing for relaxation methods
  - compute the sum of the squares of the error in each equation
  - do this at the end of each iterative loop over  $n$ 
    - stop if this is small enough



11



11

## Pseudocode for Residual Calculation

```
loop over n = 1, 2, ...
  update:      ...
  copy back: ...

  compute residue: rnorm = 0.0
                  loop over i = 1, 2, ..., M
                    rnorm = rnorm + (-u(i-1)+2*u(i)-u(i+1))2
                  end loop over i
                  rnorm = sqrt(rnorm)

  normalise:      res = rnorm / bnorm

  if (res < tolerance) finish

end loop over n
```



12



12

## Notes on Residual

- For a perfect solution, residue will be zero
  - in practice we will get a finite value
  - usually stop when it is “small”, e.g. a tolerance of  $\text{res} < 10^{-6}$
  - there will be a limit to how small the residual can get
    - can easily hit the limits of single precision
    - use double precision everywhere (or at least perform residual calculation using doubles)
- Normalisation
  - need to divide by the norm of the  $b$  vector
  - we saw before that  $b$  corresponds to the boundary values
  - in 1D:  $\text{bnorm} = \sqrt{u(0) * u(0) + u(M+1) * u(M+1)}$ 
    - in 2D, need to sum values of squares of  $u_{i,j}$  over all edges




13



13

## Residual Against Iteration



- Decreases exponentially
  - with a zero initial guess for  $u$ , should equal 1.0 at iteration zero



14



14

## Parallelisation

- Very simple for Jacobi
- Decompose the problem domain regularly across processes/threads
  - for MPI we need halo regions due to  $i \pm 1, j \pm 1$  references
  - halos are 1 cell wide for 5-point stencil
  - could be wider for larger stencils
  - swap halos between neighbouring processes every iteration
- Require global sums for, e.g., residue calculation



15



15

## Relaxation Methods

- About to cover some variations on Jacobi
  - which we hope will be faster!
- How can we tell if a method will work at all?
- Necessary (but not sufficient) condition
  - if the method arrives at the correct solution it must stay there
- Is this true for Jacobi?  $u_i^{(n+1)} = 1/2 (u_{i-1}^{(n)} + u_{i+1}^{(n)})$ 
  - for solution:  $-u_{i-1}^{(n)} + 2u_i^{(n)} - u_{i+1}^{(n)} = 0$ , i.e.  $1/2 (u_{i-1}^{(n)} + u_{i+1}^{(n)}) = u_i^{(n)} = u_i^{(n+1)}$
  - so,  $u_i^{(n+1)} = u_i^{(n)}$  and we stay at the solution
    - worth checking this for other methods



16



16



## Gauss Seidel

- Why do we need both  $u_{\text{new}}$  and  $u$ ?

```
update:    loop over internal points: i = 1, 2, ... M
            unew(i) = 0.5*( u(i-1) + u(i+1) )
            end loop over i
copy back: u(1, 2, ..., M) = unew(1, 2, ..., M)
```

- Why not do the update in place?

```
update:    loop over internal points: i = 1, 2, ... M
            u(i) = 0.5*( u(i-1) + u(i+1) )
            end loop over i
```

– this is called the *Gauss-Seidel* method




17



17

## Convergence of Gauss-Seidel



- Converges twice as fast as Jacobi
  - for less work and less storage!



18



18

## Parallelisation Gauss Seidel


- Order of the update loop is now significant
  - we used normal (*lexicographic*) order: other orderings possible
- Parallelisation of Jacobi was easy
  - Just divide grid and each processor sends its boundary data to neighbouring processor (“halo-swapping”)
- Parallelisation of Gauss Seidel is harder
  - e.g. in 1D  $u_i = \frac{1}{2}(u_{i-1} + u_{i+1})$ 

$\swarrow$   
“new” was just updated

$\nwarrow$   
“old” just about to be updated
  - Updating of every point depends on the one before, which in turns depends on the one before that...

## Parallel Gauss Seidel (wavefront)

- Consider dependencies (in 2D)
  - $u_{i,j}$  depends on recently updated values  $u_{i-1,j}$  and  $u_{i,j-1}$
- In pattern below have inter-colour dependencies,
  - E.g. updates to brown elements depend on red elements
  - ...but brown elements don't depend on each other
  - E.g.  $u_{2,2}$  does not depend on  $u_{3,1}$  or vice versa
  - Brown squares can all be calculated simultaneously
    - i.e. can do the brown (or any colour) in parallel!
- Also works in 3D (more parallelism!)




## Parallel Gauss Seidel (red-black)

- Red-black order divides grid into chequerboard
  - 2 loops: update all the red squares (in parallel) first then all the black ones (in parallel)
  - new ordering removes some dependence on already updated elements
  - enables Gauss Seidel method to be parallelised
  - ordering can affect convergence (different underlying matrix)

Processor 1

Processor 2




epcc



21

## Over Relaxation

- Recall how Jacobi solution progressed



- we have increased the value of  $u_i$  by a small amount
  - but we know the real solution is even higher
- why not increase by more than suggested
  - i.e. multiply the change by some factor  $\omega > 1$

epcc

22



22

## Over-Relaxed Gauss Seidel

- Gauss-Seidel method:  $u_i = \frac{1}{2} (u_{i-1} + u_{i+1})$ 
  - i.e.  $u_i = u_i + \frac{1}{2} [(u_{i-1} - 2u_i + u_{i+1})]$
- Multiply change (in square brackets) by  $\omega$ 
  - over-relaxed update:  $u_i = u_i + \frac{1}{2} \omega [(u_{i-1} - 2u_i + u_{i+1})]$
  - or  $u_i = (1 - \omega)u_i + \frac{1}{2} \omega (u_{i-1} + u_{i+1})$
- Notes
  - original method corresponds to  $\omega = 1$
  - if we get to a solution we stay there for any value of  $\omega$
  - Theorem by Kahan:  $\omega$  has to be in range  $(0,2)$
  - Sometimes  $\omega$  known in advance, sometimes trial and error or adaptive



23



23

## Non-Linear Equations

- Relaxation methods deal directly with equations
  - doesn't matter that we cannot express them as  $Au = b$
  - equally valid for non-linear equations (e.g. fluid dynamics)
- Non-linear equations can be very unstable
  - may need to under-relax to get convergence, i.e.  $\omega < 1$




24



24


## Extending to 2 Dimensions

- Initialise
  - set boundary values (purple)
    - zero on top, bottom and left
    - hump function on right
  - zero interior (white)
- Loop over interior
  - $i = 1, 2, \dots, M$
  - $j = 1, 2, \dots, M$
  - update  $u_{i,j}$  as appropriate
- Repeat until converged
- Write results
  - include boundaries so that the solution looks nice!



## Notes (1)

- How do we convert from  $(i, j)$  to  $(x, y)$  coordinates?
  - for a domain of size  $1 \times 1$ :
    - $x = ih$  and  $y = jh$
- What is the hump function?
  - $u(1.0, y) = k(y_2 - y)^2(y - y_1)^2$
  - a peak, centred at  $(y_2 + y_1)/2$ , dropping to zero for  $y < y_1$  and  $y > y_2$
  - for this example, take  $y_1 = 0.4$  and  $y_2 = 0.6$




## Notes (2)

- How do we convert from  $(x, y)$  to  $(i, j)$  coordinates?

- e.g. what lattice point do we look at to find  $u(0.20, 0.33)$ ?
- $(0.20, 0.33)$  is unlikely to fall exactly on a lattice point

- the four nearest neighbours are:

- $i = \text{int}(x/h)$
- $j = \text{int}(y/h)$



- do weighted average of these four values (see exercise notes)



27



27

## Convection-Diffusion Equations

- 1D Gauss-Seidel update

$$u_i = \left( \frac{1}{2 + ah} \right) (u_{i-1} + (1 + ah)u_{i+1})$$

- 1D Over-Relaxed update

$$u_i = (1 - \omega)u_i + \omega \left( \frac{1}{2 + ah} \right) (u_{i-1} + (1 + ah)u_{i+1})$$

- 2D Discrete Equations

$$u_{i,j} = \frac{1}{(4 + (a_x + a_y)h)} (u_{i,j-1} + u_{i-1,j} + (1 + a_x h)u_{i+1,j} + (1 + a_y h)u_{i,j+1})$$

$(a_x, a_y)$  = wind strength from  $x$  (East) and  $y$  (North) respectively



28



28

## Notes

- Have multiplied all the equations by  $h^2$ 
  - equations now explicitly depend on  $h$  for a non-zero wind  $a$
  - straightforward to derive update equations for 2D case
- A different convention for Krylov methods (later)
  - maintain the  $1/h^2$  factor in matrix  $A$ 
    - therefore need to multiply RHS by same factor
    - happens to be more convenient
- Finite wind
  - matrix  $A$  is now non-symmetric
  - in 1D, lower-diagonal elements are  $(1 + ah)$ , upper elements are 1
  - gives some minor technical issues when normalising the residue
    - see practical notes
    - if correctly normalised, residue at zero iterations will *always* be 1.0 if the initial guess is a zero solution



29



29

## Summary

- Relaxation methods
  - guess at an initial solution
  - update many times and stop when residue is small enough
- Update rule is very straightforward
  - solve exactly for each individual  $u_i$ 
    - obtain formula by rearranging difference equations so  $u_i$  is on the LHS
- Interior points updated according to the PDE
  - boundary points set by the boundary conditions
- Jacobi is the simplest method
  - Gauss Seidel acts “in-place” and requires roughly half the iterations
  - appropriate over-relaxation can accelerate this even more
    - finding the best value of  $\omega$  requires some experimentation!



30



30