

CSR format with empty row

$$A = \begin{pmatrix} 1.0 & 0.1 & 0.5 & \cdot \\ 0.4 & \cdot & \cdot & 0.7 \\ \cdot & \cdot & \cdot & \cdot \\ 0.3 & \cdot & 0.2 & 0.8 \end{pmatrix}$$

```
nRow = 4; nCol = 4; nzMax = 8;  
value[8] = {1.0, 0.1, 0.5, 0.4, 0.7, 0.3,  
0.2, 0.8};  
colIdx[8] = {0, 1, 2, 0, 3, 0, 2, 3};  
rowStart[5] = {0, 3, 5, 5, 8};
```



1



1

Numerical Algorithms for HPC

KS methods in parallel, Preconditioning, matrix splitting
and KS methods in action



2

Outline

- KS method in parallel
 - Matrix-vector multiplication
- Preconditioning
 - Basic concept
 - Matrix splitting
- CG and BiCGStab in action



3



3

Matrix vector multiplication in parallel

- Krylov Subspace methods involve one or more matrix-vector multiplications at each iteration
- The matrix-vector multiplication is the most computationally expensive part of most KS methods
- It makes sense to carry this out in parallel where possible
- We first look at trivially parallel case of block-diagonal matrix
- Then look at more general case for
 - Replicated vector
 - Distributed vector



4



4

Parallel matrix vector – block diagonal matrix

- How to decompose matrix across parallel machine?
 - Depends on how the matrix is stored
 - Depends on what the matrix is
- Block diagonal matrix-vector multiply routine is completely parallelisable

$$\begin{array}{c} P_1 \\ P_2 \\ P_3 \end{array} \begin{array}{|c|c|c|} \hline A_1 & & \\ \hline & A_2 & \\ \hline & & A_3 \\ \hline \end{array} \begin{array}{c} v_1 \\ v_2 \\ v_3 \end{array} = \begin{array}{c} A_{11} \cdot v_1 \\ A_{22} \cdot v_2 \\ A_{33} \cdot v_3 \end{array} \begin{array}{c} P_1 \\ P_2 \\ P_3 \end{array}$$

- Could further parallelise via threads within blocks



5



5

General parallel matrix-vector

- Multiplication of $Av = w$
- General approach is 1D decomposition of matrix A by row
 - each processor has $m = N_{rows}/N_{procs}$ rows of matrix
- Then have a choice as to whether the vector v is
 - Replicated across all processors – all processors can see all the entire vector at each iteration
 - Distributed across processors – processors only see subset of data at any one time




6



6

Parallel matrix-vector – Replicated data

- Consider first processor, P_1



- All processors have copy of entire vector, v
- Need to broadcast data at end of each iteration to complete the vector w
- Vector w then replicated across procs and ready for next iteration
- Simple to implement with no communication during calculation
 - but high memory requirement and requires broadcast at the end

epcc

7



7

Distributed vector: 2 processor 2x2 example

Both matrix **and** vector are distributed:

$$\begin{array}{l} P_1 \\ P_2 \end{array} \left(\begin{array}{cc} a_{11} & a_{12} \\ a_{21} & a_{22} \end{array} \right) \left(\begin{array}{c} b_1 \\ b_2 \end{array} \right) = \left(\begin{array}{c} a_{11}b_1 + a_{12}\textcircled{b_2} \\ a_{21}\textcircled{b_1} + a_{22}b_2 \end{array} \right)$$

- Circled elements not initially available to relevant processor
- Calculate $a_{11}b_1$ on P_1 and $a_{22}b_2$ on P_2
- Then swap elements b_1 and b_2 between processors
- Then calculate $a_{12}b_2$ on P_1 and $a_{21}b_1$ on P_2
- Can generalise to more processors
 - e.g. pass around a ring rather than simply “swap” elements
- Can also generalise to bigger matrices
 - Then deal with blocks of rows instead
- No need to gather at the end - vectors remain distributed!

epcc

8



8

Parallel matrix-vector – Distributed data

- Distributed vector
 - more complicated than replicated and more communication during calculations
 - but more efficient use of memory for storing both matrix **and** vectors
 - answer ends up distributed as before but can stay distributed!
 - No need for final broadcast of data
 - locally accumulate matrix times vector block (i.e. partial dot products)
 - processors must exchange vector blocks to complete dot products
 - Sparse: may not need all processors to exchange blocks
- Smart implementations for specific matrix types
- Global sum still required if scalar products needed (e.g. residues)



9



9

Preconditioning: Overview

- Motivation
- What is preconditioning?
- What is its purpose?
- Common preconditioners



10



10

What is preconditioning?

“a preconditioner is any form of implicit or explicit modification of an original linear system which makes it “easier” to solve by a given iterative method”

Y Saad, Iterative methods for Sparse Linear Systems

- Examples of preconditioners
 - Scaling all rows so that diagonal entries are equal to 1 (Jacobi preconditioner)
 - See block Jacobi a few slides later...
 - Pre-multiplying the matrix by a given matrix, e.g. $A \rightarrow M^{-1}A$
 - Unlikely M or $M^{-1}A$ ever computed directly
 - M^{-1} may be complicated. E.g. result of some FFT transformations or integral calculations



11



11

What is preconditioning? cont ...

- Two extremal cases:
 - Choice $M = I$ is equivalent to no preconditioning
 - choice $M = A$ is equivalent to factorising the problem directly
- We seek an intermediate M which ensures that KS method will converge and reduces (minimises) the cost of solver.
- Idea to preserve structure of A (particularly symmetry!)



12



12

Motivation

- We would like Krylov subspace method
 - to converge (smoothly)
 - to converge in as few iterations as possible
 - reduce the effects of rounding error
 - make method more tractable for multiple righthand sides
- Number of iterations is affected by condition number, C , of the matrix
 - $C = \frac{\lambda_{\max}}{\lambda_{\min}}$, where λ_{\max} and λ_{\min} are the maximum and minimum eigenvalues of the matrix respectively
 - a lower condition number implies fewer KS iterations
- Choose preconditioner which reduces the condition number



13



13

Preconditioning in linear systems

- Solving system

$$Au = b$$

- is equivalent to solving

$$M^{-1}Au = M^{-1}b$$

- where M is a SPD matrix (for CG, at least).
- Idea is to choose M such that similar to A but easier to invert.
- Jacobi/Gauss-Seidel can be thought of as pre-conditioner (see matrix splitting later)




14



14

Block Jacobi preconditioning

- Ignore entries that are not in block diagonals
- Invert blocks locally (often incomplete inverse)



- Blocks get smaller with more processors so less is inverted

epcc

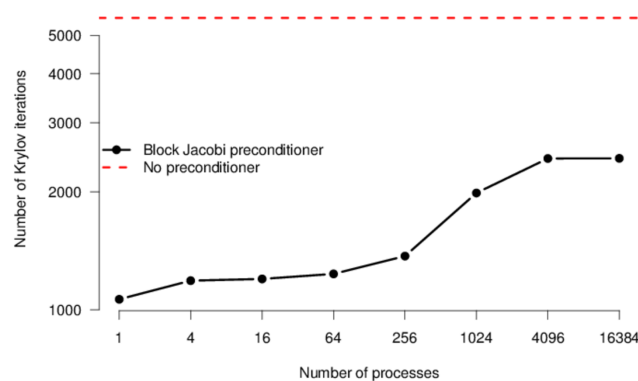
15



15

More processors give worse approximation

- Throws away more information



epcc

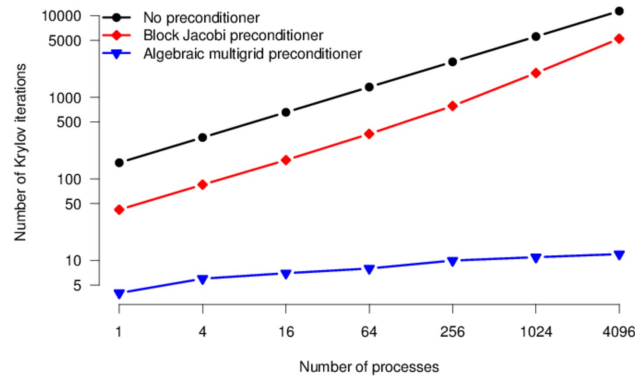
16



16

Other preconditioners exist

- Multigrid (algorithmically optimal)
 - computationally more difficult and more expensive



Incomplete LU factorisation

- A more powerful/expensive preconditioner is ILU factorisation
- Elements are computed as in LU factorisation, but those that fall outwith the sparsity pattern are discarded

$$A = LU - R$$

- $M = LU$ can then be used as the preconditioner
- This preserves sparsity pattern
- No guarantee of existence of non-singular ILU factors.
- Many modifications exist

Preconditioned CG

- Remember, we were originally solving the linear system

$$Av = b$$

- With the preconditioner, we are now solving the equally valid linear system

$$M^{-1}Av = M^{-1}b$$

- Remember the original definition of the residual

$$r = b - Av$$

which would become, with preconditioning

$$M^{-1}r = M^{-1}(b - Av)$$

- We can introduce a preconditioned analogue, s :

$$s = M^{-1}r \Rightarrow Ms = r$$



19



19

Preconditioned CG algorithm

Choose v_0 , compute $r_0 = b - Av_0$, $k=0$,

Solve $Ms_0 = r_0$ (using a direct method), $p_0 = s_0$

While ($k < \text{maxiter}$)

$$\alpha = r_k \cdot s_k / p_k \cdot A p_k$$

$$v_{k+1} = v_k + \alpha p_k$$

$$r_{k+1} = r_k - \alpha A p_k$$

if ($\|r_{k+1}\|_2 / \|b\|_2 < \text{tol}$) break

Solve $Ms_{k+1} = r_{k+1}$

$$\beta = r_{k+1} \cdot s_{k+1} / r_k \cdot s_k$$

$$p_{k+1} = s_{k+1} + \beta p_k$$

$k = k + 1$

end while



20



20

Matrix Splitting

- Returning to Jacobi and Gauss Seidel...
 - Any linear problem is of the form $Ax = b$
 - A encodes the precise form of the PDE
 - b contains any fixed boundary conditions
 - Could split A into three parts
 - Diagonal, Strictly Upper and Strictly Lower triangular: $A = L + D + U$
 - *not* the same as the LU factors!
- $$(L + D + U)x = b, \quad Dx = -(L + U)x + b$$
- view these as iterative expressions, e.g. Jacobi corresponds to
- $$Dx^{(n+1)} = -(L + U)x^{(n)} + b$$
- $$x^{(n+1)} = -D^{-1}(L + U)x^{(n)} + D^{-1}b$$

Consider 1D Pollution Problem

- A represents: $-\frac{d^2}{dx^2}$

$$A = \begin{bmatrix} 2 & -1 & \cdot & \cdot \\ -1 & 2 & -1 & \cdot \\ \cdot & -1 & 2 & -1 \\ \cdot & \cdot & -1 & 2 \end{bmatrix}$$

- Splitting into L, D and U

– Jacobi iteration is actually given by: $Dx^{(n+1)} = -(L + U)x^{(n)} + b$

$$\begin{bmatrix} 2 & \cdot & \cdot & \cdot \\ \cdot & 2 & \cdot & \cdot \\ \cdot & \cdot & 2 & \cdot \\ \cdot & \cdot & \cdot & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n+1)} = \begin{bmatrix} \cdot & 1 & \cdot & \cdot \\ 1 & \cdot & 1 & \cdot \\ \cdot & 1 & \cdot & 1 \\ \cdot & \cdot & 1 & \cdot \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n)} + \begin{bmatrix} b_1 \\ 0 \\ 0 \\ b_4 \end{bmatrix}$$

Jacobi Equations

- Equations the same as in previous lectures
 - with u_i replaced by x_i
 - exterior boundary values u_0 and u_{N+1} replaced by b_1 and b_N

$$\begin{aligned} x_1^{(n+1)} &= \frac{1}{2} (b_1 + x_2^{(n)}) \\ x_2^{(n+1)} &= \frac{1}{2} (x_1^{(n)} + x_3^{(n)}) \\ x_3^{(n+1)} &= \frac{1}{2} (x_2^{(n)} + x_4^{(n)}) \\ x_4^{(n+1)} &= \frac{1}{2} (x_3^{(n)} + b_4) \end{aligned}$$

- Procedure
 - impose PDE at each interior point
 - new value is the average of the old neighbouring points



23



23

Gauss Seidel

- Keep both \mathbf{D} and \mathbf{L} on the LHS: $(\mathbf{D} + \mathbf{L})\mathbf{x}^{(n+1)} = -\mathbf{U}\mathbf{x}^{(n)} + \mathbf{b}$

$$\begin{bmatrix} 2 & \cdot & \cdot & \cdot \\ -1 & 2 & \cdot & \cdot \\ \cdot & -1 & 2 & \cdot \\ \cdot & \cdot & -1 & 2 \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n+1)} = \begin{bmatrix} \cdot & 1 & \cdot & \cdot \\ \cdot & \cdot & 1 & \cdot \\ \cdot & \cdot & \cdot & 1 \\ \cdot & \cdot & \cdot & \cdot \end{bmatrix} \begin{bmatrix} x_1 \\ x_2 \\ x_3 \\ x_4 \end{bmatrix}^{(n)} + \begin{bmatrix} b_1 \\ 0 \\ 0 \\ b_4 \end{bmatrix}$$

$$\begin{aligned} x_1^{(n+1)} &= \frac{1}{2} (b_1 + x_2^{(n)}) \\ x_2^{(n+1)} &= \frac{1}{2} (x_1^{(n+1)} + x_3^{(n)}) \\ x_3^{(n+1)} &= \frac{1}{2} (x_2^{(n+1)} + x_4^{(n)}) \\ x_4^{(n+1)} &= \frac{1}{2} (x_3^{(n+1)} + b_4) \end{aligned}$$

- equivalent to solving Jacobi equations *in-place* in order 1, 2, ..., N



24



24

Jacobi and (over-relaxed) Gauss-Seidel

- Connection to matrix-splitting

- Jacobi

$$\mathbf{x}^{(n+1)} = -\mathbf{D}^{-1}(\mathbf{L} + \mathbf{U})\mathbf{x}^{(n)} + \mathbf{D}^{-1}\mathbf{b}$$

- Gauss-Seidel

$$\mathbf{x}^{(n+1)} = -(\mathbf{D} + \mathbf{L})^{-1}\mathbf{U}\mathbf{x}^{(n)} + \mathbf{D}^{-1}\mathbf{b}$$

- Over-relaxed Gauss-Seidel

$$\mathbf{x}^{(n+1)} = (\mathbf{D} + \omega\mathbf{L})^{-1} \left(((1 - \omega)\mathbf{D} - \omega\mathbf{U})\mathbf{x}^{(n)} + \omega\mathbf{b} \right)$$

- All have the form

$$\begin{aligned}\mathbf{x}^{(n+1)} &= \mathbf{E}^{-1}\mathbf{F}\mathbf{x}^{(n)} + \mathbf{E}^{-1}\mathbf{b} \\ \mathbf{E}\mathbf{x}^{(n+1)} &= \mathbf{F}\mathbf{x}^{(n)} + \mathbf{b}\end{aligned}$$



25



25

Note: Convergence of Splitting

- General matrix splitting eqns: $\mathbf{E}\mathbf{x}^{(n+1)} = \mathbf{F}\mathbf{x}^{(n)} + \mathbf{b}$
 - solution at iteration n is perfect solution $\hat{\mathbf{x}}$ plus correction $\delta\mathbf{x}^{(n)}$

$$\mathbf{x}^{(n)} = \hat{\mathbf{x}} + \delta\mathbf{x}^{(n)} \text{ where } \mathbf{E}\hat{\mathbf{x}} = \mathbf{F}\hat{\mathbf{x}} + \mathbf{b}$$
 - substituting into main equation gives $\mathbf{E}\delta\mathbf{x}^{(n+1)} = \mathbf{F}\delta\mathbf{x}^{(n)}$
 - error in solution evolves according to $\delta\mathbf{x}^{(n)} = (\mathbf{E}^{-1}\mathbf{F})^n \delta\mathbf{x}^{(0)}$
- Convergence depends on eigenvalues of $\mathbf{E}^{-1}\mathbf{F}$
 - must all be less than one in order to get a solution
 - speed of convergence depends on condition number
- Iteration matrix
 - $\mathbf{E}^{-1}\mathbf{F}$ is $-\mathbf{D}^{-1}(\mathbf{L} + \mathbf{U})$ for Jacobi and $-(\mathbf{D} + \mathbf{L})^{-1}\mathbf{U}$ for Gauss Seidel
 - can show that latter is better conditioned
 - heuristically, GS inverts more of the matrix at each step



26



26

CG and BiCGstab in action

- Next few slides look at KS in action for pollution problem
- Consider asymmetry (wind)
- Look at convergence




Asymmetry - wind

- With wind the matrix A is
 - not symmetric
 - positive definite
- CG properties no longer guaranteed
- Gentle breeze,
 - CG still works
 - Takes longer to converge
- Increase wind until CG breaks down
 - What happens to the norm of the residual?
- Implement BiCGstab
 - Play with a hurricane!

A still day

$$\text{Wind} = F (1.0 x + 0.5 y)$$

$$F = 0.0$$



epcc



29


A gentle breeze

$$\text{Wind} = F (1.0 x + 0.5 y)$$

$$F = 1.0$$

Change in residual no longer monotonic

Algorithm is no longer progressive



epcc




30

A moderate breeze

$$\text{Wind} = F (1.0 x + 0.5 y)$$

$$F = 1.5$$

Number of iterations required to reach residual increases as wind strength increases



epcc



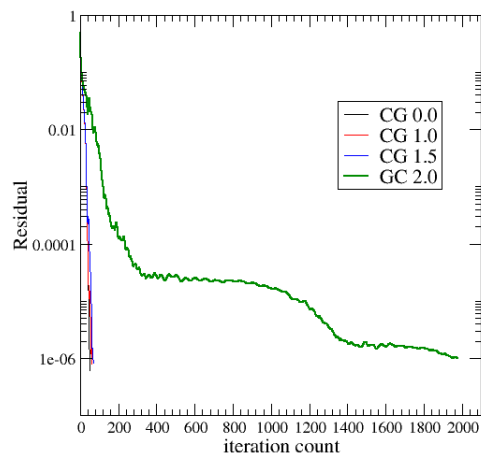
31

A stiff wind

$$\text{Wind} = F (1.0 x + 0.5 y)$$

$$F = 2.0$$

Matrix further from SPD
CG is starting to break down
Now takes thousands of iterations



epcc



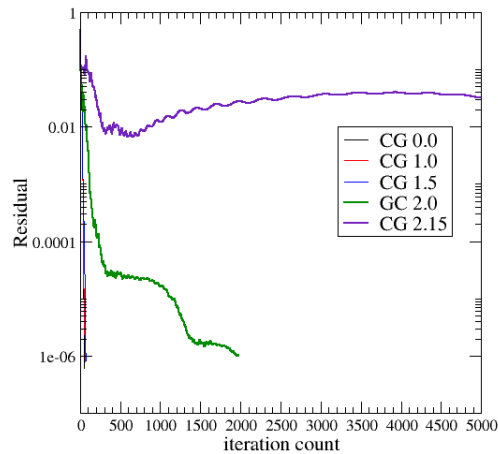
32

Slightly stronger wind

Wind = $F (1.0 x + 0.5 y)$

$F = 2.15$

CG fails to converge!



epcc



33


BiCGstab

Wind = $F (1.0 x + 0.5 y)$

BiCGstab converges in few iterations for even large winds

Large winds off diagonal terms dominate

BiCGstab does **twice** as much work per iteration as CG



epcc



34

Conclusions

- Matrix-vector multiplication in parallel with vectors stored as
 - replicated data
 - distributed data
- Preconditioning of linear system can improve reliability of KS method, reduce iteration count and computational costs
- Simple preconditioners based on stationary splitting methods
- More complex methods such as ILU, or preconditioning using Fourier transforms are more effective but may not work at all.
- Also looked at CG and BiCGstab in action



35



35

Remaining exercises

- Pollution model
 - Write a method that converts a COO-formatted matrix into a CSR-formatted matrix
 - Write a method that performs a matrix-vector multiplication for a CSR-formatted matrix.
 - Implement CG
- Matrix-vector in parallel – no practical session for this
 - Take serial “power method” eigensolver and parallelise the matrix-vector routine for both replicated data and distributed data



36



36