

Numerical Algorithms for HPC

FFT Libraries



1



1

Overview

- Recap of cost of DFT and FFT
- Use of Libraries
- FFT Libraries
 - General usage notes
- FFTW
 - Overview
 - Using the Library
 - Performance
- FFT Practical



2



2

Recap: Naïve DFT cost

- DFT is a sum. For each of the N values of F_n there is a sum of length N to perform

$$F_n = \sum_{k=0}^{N-1} f_k e^{2\pi i k n / N}$$

- Cost is therefore $N \times N = N^2$ i.e. $\mathcal{O}(N^2)$



3



3

Recap: FFT cost

- FFT split sum into 2, then into 2 again, etc. E.g. for $N = 4$

f_0	f_2	f_1	f_3	Bit reversal: $\mathcal{O}(N)$
$f_0 + f_2$	$f_0 - f_2$	$f_1 + f_3$	$f_1 - f_3$	
$f_0 + f_2 + f_1 + f_3$	$f_0 - f_2 + if_1 - if_3$	$f_0 + f_2 - f_1 - f_3$	$f_0 - f_2 - if_1 + if_3$	

- $N = 4$ case: each of the $N = 4$ elements performs 2 sums
- The number of times N (if a power of 2) can be halved until you get to 1 is $\log_2 N$
 - E.g. for $N=32$: $32/2=16$, $16/2=8$, $8/2=4$, $4/2=2$, $2/1=1$
 - Divides 5 times: $\log_2 32 = 5$ (i.e. $2^5 = 32$)
- So in general cost is $\mathcal{O}(N \log_2 N)$



4



4

Use of Libraries

- What is a library?
 - Collection of pre-implemented routines available in some sort of package (static archive, shared library, Framework, DLL...)
 - Defined (and hopefully documented) API
 - Often provided with operating system (e.g. Scientific Linux), available via package managers, or binary/source downloads
 - Can be combined with your code by linking (and perhaps #including header files)



5



5

Use of Libraries

- Why use libraries?
 - Easier - save effort and time, no duplicated code
 - Avoid bugs – libraries are (hopefully) well tested
 - Performance – usually contain efficient implementations, perhaps optimised for a given system (vendor-specific architectures), or auto-tuning



6



6

FFT Libraries

- FFTs do not normalise
 - Each FFT/Inverse FFT pair scales by a factor of N
 - Usually left as an exercise for the programmer.
- DFTs are complex-to-complex transforms, however, most applications require real-to-complex transforms
 - Simple solution: set imaginary part of input data to be zero
 - This will be relatively slow
 - May be better to pack and unpack data
 - Place all the real data into all slots of the input, complex array (of length $N/2$) and then unpack the result on the other side ($O(N)$)
 - Around twice as fast as the simple solution
 - Good details in Numerical Recipes book
 - Some libraries have real-to-complex wrappers
 - Similar approach allows us to do 2 independent real-to-complex transforms simultaneously



7



7

FFT Libraries

- Multidimensional FFTs
 - simply successive FFTs over each dimension
 - order immaterial: linearly independent operations
 - can place data into 1D array
 - strided FFTs
 - some libraries have multidimensional FFT wrappers
- Parallel FFTs
 - performing FFT on distributed data
 - 1D FFTs are cumbersome to parallelise
 - Suitable only for huge N
 - 2+D parallel, array transpose operation
 - distributed data is collated on one processor before FFT
 - more in parallel libraries lecture later
 - some libraries have parallel FFT wrappers



8



8

FFTW

- Fastest Fourier Transform in the West
 - www.fftw.org
- The FFTW package was developed at MIT by Matteo Frigo and Steven G. Johnson
- Free under GNU General Public License
- Portable, self-optimising C code
 - Runs on a wide range of platforms
- Arbitrary sized FFTs of one or more dimensions
 - Fastest routines where N is composed of powers of 2, 3, 5 and 7 (other sizes can be optimised for at configuration time)
 - Also 11 and 13 if to power 0 or 1
 - Powers of 2 particularly fast!



9



9

FFTW: FFTW2 vs FFTW3

- Previous version: “FFTW2”
 - Many legacy codes employ FFTW2
 - Simple(r) C interface, with wrappers for many other languages
 - Supports MPI
- New version: “FFTW3”
 - Different interface to FFTW2, though not radically – to allow planner more freedom to optimise
 - Users must rewrite code
 - Now has MPI support
 - Somewhat faster than FFTW2 (~10% or more)
 - Rest of this lecture assumes “FFTW3”



10



10

FFTW: Details

- Can perform FFTs on distributed data
 - MPI for distributed memory platforms
 - OpenMP or POSIX for SMPs
- If users rewrite their code to use FFTW just once then the user is saved from
 - learning platform dependent, proprietary FFT routines
 - rewriting their code every time they port their code
 - No standard interface to FFTs
 - drastically rewriting their makefiles
 - location of FFTW libraries may vary
- Fortran API is available, though modern best practice is to use bindings to C.
- FFTs can not always be done “in-place”
 - The input and output arrays must be separate and distinct or temp arrays created



11



11

FFTW: Plans

- All FFT libraries pre-compute the *twiddle factors*
- FFTW ‘plans’ also generate the FFT code from *codelets*, hard-coded small transforms
 - Codelets compiled when FFTW configured
- Two forms of plans
 - Estimated
 - The best numerical routines are guessed, based on information gleaned from the configuration process.
 - Measured
 - Different numerical routines are actually run and timed with the fastest being used for all future FFTW calls using this plan.
 - Variants widen or restrict range of algorithms to consider.
- Old plans can be reused or even read from file: *wisdom*



12



12

FFTW: Installation

- Download library from the website and unpack (gzipped tar file)
- `./configure; make; make install`
 - Probes the local environment
 - Compiles the many *codelets* (small C object codes)
 - User can provide non-standard compiler optimisation flags
 - Libraries (both static and dynamic) are then installed along with online documentation and header files
- Includes test suite
 - Very important for any numerical library
- Pre-installed with Scientific Linux
- You won't have to install FFTW for exercise – it's already available on Cirrus via a module:
`module load fftw/3.3.8-gcc8-mpi4`



13



13

FFTW: Compiling code

- `gfortran fft_code.f -O3 -lfftw3`
 - If using C, FFTW must be linked with `-lfftw3 -lm`
 - Different precision libraries may be available (`fftw3f` and `fftw3l` for float and long double respectively).
 - May need to include location of libraries (`-L/...`) or include files (`-I/...`).
- Example Fortran code using C bindings:

```
type(c_ptr) :: plan
integer(c_int), parameter :: n = 1024
complex(c_double_complex) :: in(n), out(n)
! plan the computation
plan = fftw_plan_dft_1d(...)
! execute the plan
call fftw_execute_dft(...)
```

 - NB: full incantations are not given here as reading documentation is integral to utilising any numerical library



14



14

FFTW: Performance

- The FFTW homepage, www.fftw.org, details the performance of the library compared to proprietary FFTs on a wide range of platforms.
- The FFTW library is faster than any other portable FFT library
- Comparable with machine-specific libraries provided by vendors
- Performance results from <http://www.fftw.org/speed/>




15



15

FFTW: AMD Opteron 275 2.2 GHz




16



16

FFTW: Intel Core Duo 3.0 GHz




epcc

17



17

FFTW: IBM POWER5 1.65 GHz



epcc

18



18

FFTW: Accolades

- Winner of the 1998 J.H. Wilkinson Prize for Numerical Software
 - awarded every four years to the software that "best addresses all phases of the preparation of high quality numerical software."
- Quotes from www.fftw.org.
 - "Its the best FFT package I have ever seen"
 - "It performs my standard 256 iterations of 1024pt complex FFT about 20 times faster than the previous one I used."
 - "FFTW is the best thing since microwave popcorn"
- Dr Richard Field
 - Former Vice-principal of University of Edinburgh and Chairman of NAG
 - "I think FFTW is terrific. Its the best piece of software I've seen written in a bunch of years. [...] I give FFTW very high marks (probably as high marks as I would ever give)."



19



19

FFT Summary

- First lecture:
 - Introduced both the continuous and discrete forms of the Fourier Transform
 - Fast Fourier Transform
- Second lecture:
 - FFTW
 - Fast, robust and portable
 - Fortran and C, serial and parallel.
 - Simple to use
 - Recommended and used in major projects by EPCC



20



20

FFT Practical

- FFTs are also employed in JPEG image compression and smoothing...
 - MPEG also utilises FFTs along with other graphical techniques
- Examine how manipulating image data in Fourier space alters the transformed image using FFTW 3



epcc

21



21

FFT Practical

- Full instructions in practical handout sheet
- Two main image transformations to attempt:
 - How many higher frequencies can be removed, reducing the size of the image in memory, without visibly altering the image?
 - Performing edge detection by removing only lower frequencies.
- Hint: remember that Fourier space is periodic
- Begin with an estimated plan and a 1D, non-strided, complex-to-complex transform

epcc

22



22

FFT Practical

- Fortran and C makefiles and templates available
- For exercises
 - Fortran routines with C bindings:
 - `fftw_plan_dft_1d`, `fftw_execute_dft`, `fftw_destroy_plan`
 - C routines:
 - `fftw_plan_dft_1d`, `fftw_execute`, `fftw_destroy_plan`
 - NB: Read the documentation for arguments, usage etc.
- Online documentation on FFTW
 - the `info fftw` command or www.fftw.org



23



23

FFT Practical : Further work

- How competitive is FFTW compared to other FFT libraries?
 - Alter previous code and time sections for plans and FFTs for FFTW and FFTPACK, say.
 - At present, FFTW is the only FFT library available
 - If interested, install FFTPACK in your home space
 - Extend data array to compare FFT libraries over a range of N
 - Including non-power-of-two N
 - fill array with zeros first
- Convert image manipulation code to use
 - 2D FFTW
 - Investigate the use of “wisdom”



24



24