

Numerical Algorithms for HPC

Random Number Generation

Introduction

- Random numbers are frequently used in many types of computer simulation
- Frequently as part of a sampling process:
 - Generate a representative sample of a large population by choosing members at random.
 - Monte-carlo integration is approximating an integral by sampling the function at random points.
 - Even when simulating a stochastic process (random walk/random events etc.) we are sampling the possible evolutions of the system.

What is Random anyway

- “Random” is actually a very difficult philosophical concept.
- However in most cases the real requirement is “unbiased sampling” which is more straightforward.

Distribution

- Random numbers are chosen from a probability distribution.
- For random integers each possible result X occurs with a probability $P(X)$
- For random real numbers R this becomes a probability density $P(R)$
 - Chance of the results occurring within a region is the integral of the probability density over that region.
 - Most generators are designed to generate a “uniform” distribution.
 - $P(R) = 1 \quad 0 \leq R < 1$
 - $P(R) = 0 \quad \text{elsewhere}$
 - Other distributions then generated from this

Resolution

- However computers use floating point numbers not true real numbers.
 - Only a finite set of possible values can be represented.
 - Any random “real” number must come from this set.
- Most techniques generate an even smaller sub-set of values e.g.
 - $R = X/N$ where X is a random integer between 0 and $N-1$
 - $1/N$ is the resolution of that generator.
- Generated distribution is only an approximation to uniform.
 - May bias the results if you are not careful
 - Always worth understanding the resolution of your generator

Correlations

- True Random numbers are also un-correlated with each other.
- The probability of getting a particular set of random results should be the product of the probabilities of each result in isolation.

Hardware Random Number Generators

- You can build hardware random number generators.
 - These work by taking measurements of some random physical process
 - Thermal noise
 - Quantum processes.
 - Problems
 - Debugging is very hard as can never reproduce the same program run twice.
 - May still suffer from limited resolution
 - Often quite slow.
 - May not result in any visible improvement in quality of results.
 - More commonly used in cryptographic applications.

Pseudo Random Numbers

- Pseudo Random Numbers are a deterministic sequence of numbers generated by some algorithm that are used in-place of true random numbers.
- Aim is for the sequence to share enough of the statistical properties of true random numbers not to bias the results.
- PRNs are NOT random. It is always possible to come up with some test that demonstrates this.

PRNG Quality

- Quality of a PRNG sequence depends on the intended use.
 - Each use case only depends on *some* of the statistical properties of true random numbers.
 - Some generators may introduce problems for some calculations but not others.
- In practice, algorithms exist that can stand in for true random numbers for most common types of simulation.
- Unfortunately language default generators are often fairly poor.

Structure of a PRNG

- Logically PRNGs consist of:
 - An internal state S_i
 - An update transform $T S_i \rightarrow S_{i+1}$ that maps one state onto the next
 - An output transform $F S_i \rightarrow X_i$ that generates the next number in the PRNG sequence from the current state.
- Algorithms are rated on the statistical properties of the output sequence
 - Speed of execution and memory consumption are also important.
- Different algorithms may generate the same PRNG sequence via different state representations and transforms.

Seeding the Generator

- Also need some mechanism of initialising the starting state.
- Traditional algorithms only used a single word of state so many programs assume the generator is initialised using a single integer.
- If you don't set a starting seed you either:
 1. Get the same sequence every time you run the program.
 2. The generator seeds from the current time (makes debugging hard).
- If your program checkpoints remember to save RNG state so you can restart **exactly** where you left off.
 - Write tests to check this!

Period of a generator

- There are only a finite number of possible states.
- Eventually generator will return to its starting state.
- The update transform should generate a cyclic group
 $T^{period} = I$
- The size of this group is the *period* of the generator.
- It is also the number of valid states.

How state is stored

- In principle you could store the position in the sequence.
 - Update transform is just an increment $i \rightarrow i + 1$
 - All the randomness is in the output-transform.
 - Need very expensive output-transform to have good randomness properties.
- In practice use state representations that approximate random values and keep the output transform simple.
 - Even fairly simple (inexpensive) update transforms can have good randomness properties

PRNG Algorithms

- PRNG Algorithms are deceptively simple.
 - Usually made up from a few simple operations.
 - Typically bitwise operations or modular arithmetic.
- Very tempting to try and “Improve” on published algorithms
- **DON'T DO THIS** unless you really know what you are doing.
- Each new algorithm requires theoretical (Number theory) analysis to determine the period of the generator.
 - Many other statistical properties can also be derived theoretically.

Selecting Generators

- Most generators are selected based on the properties of small sets of consecutive numbers from the sequence.
 - $\{X_i, X_{i+1}\}$ approximate a pair of random number.
- Non consecutive sets may appear less random.
 - E.g $\{X_i, X_{i+1024}\}$
- Consecutive sets important for most applications (especially when used to generate non-uniform distributions) so this is a good heuristic for general purpose generators.
- For a specific application may be other correlations that are equally important.

- Selection uses a combination of theory and statistical tests.
- Statistical tests augment theory, not good enough by themselves.

Linear Congruential Generators

- $S_{i+1} = (a S_i + c) \bmod M$
- If a , c and M chosen correctly, has M possible states.
- If $c = 0$ then $(M-1)$ possible states ($S=0$ always maps to itself).

Java.util.Random

- Optimised for speed not quality
 - $a = 0x5DEECE66DL$
 - $C = 0xBL$
 - $M = 2^{48}$
- $\text{Mod } 2^{48}$ is a bit-mask so very fast.
- 47 bits of state in total.
 - However bit- n of the state has repeats with at most period 2^{n+1}
 - bit-0 period 2
 - bit-1 period 4
 - Only the high order bits repeat with any degree of randomness.
 - Class only exposes the top 32-bits to the user making it ok for quick and dirty use.

MRGs

- LCG are a special case of Multiply Recursive Generators
 - $S_n = a_1 S_{n-1} + \dots + a_k S_{n-k} \bmod M$
 - Needs array of state variables.
- Some number theory ...
 - If $M = P^q$ with P prime then maximum period is $P^{q-1}(P^k - 1)$.
 - Special values of $\{a_k\}$ generate full period if M is prime.
- Many other common generators are special cases of these.

Other Common generators


- LFSR
 - Linear Feedback Shift Register
 - $M=2$
- Lagged Fibonacci Generators
 - Only 2 Values of $\{a_k\}$ non-zero so faster than the general case.
- Mersenne-Twister appears quite different but is equivalent to a MRG with $M=2$ and $(2^k - 1)$ a Mersenne prime.

Non uniform distributions

- Non-uniform distributions are constructed out of (multiple) normally distributed values.
- For any probability distribution $p(x)$
 - $\int_{min}^{max} p(x) = 1$
 - Selecting small areas under the curve uniformly is the same as selecting x with probability $p(x)$
 - Inverse transform sampling.
 - Divide area into thin strips of equal area and select strip at random.
 - Rejection sampling
 - Choose x, y points at random but reject points above curve i.e. $y > p(x)$

Simple example


- $p(x) = 3x^2$



Inverse transform sampling

- Generally quite hard to do:
 - Generate uniform deviate U .
 - Return x : $\int_{min}^x p(y)dy = U$
- Only analytically solvable for certain distributions.
 - e.g for $p(x) = 3x^2$
 - $x = \sqrt[3]{U}$ (cube root of U)
 - 100,000 samples & 100 bins

```
call random_number(myrng)  
myrng = myrng**(1.0/3.0)
```




Rejection sampling

- Only need to be able to evaluate $p(x)$.
 - Needs special handling for unbounded distributions.

- e.g for $p(x) = 3x^2$

```
call random_number(myrng1)
call random_number(myrng2)
myrng2 = 3.0*myrng2
if (myrng2 < 3.0*myrng1**2)
  myrng = myrng1
```




Generating Gaussians

- Most commonly required non-uniform distribution is the normal / gaussian distribution

$$- P(x) = \frac{1}{\sigma\sqrt{2\pi}} e^{\frac{-x^2}{2\sigma^2}}$$

-e.g for $\sigma = 0.5$



Box Muller

- Generates pairs of gaussians from pairs of uniform.
 - Generate 2 Uniform random numbers U, V from $(0:1]$
 - $X = \sqrt{-2 \ln U} \cdot \cos(2\pi V)$
 - $Y = \sqrt{-2 \ln U} \cdot \sin(2\pi V)$
- Generally quite slow due to math library functions.
- With care can be vectorised so may be better algorithm for GPGPU.

Polar method

- Variation of box-muller that uses accept-reject step instead of trig functions.

1. $a = 2 U - 1$

2. $b = 2 V - 1$

3. $s = a^2 + b^2$

4. *If $s > 1$ goto (1)*

5. $X = a \sqrt{\frac{-2 \ln(s)}{s}}$

6. $Y = b \sqrt{\frac{-2 \ln(s)}{s}}$

- Usually faster overall but accept/reject inhibits vectorisation

Summary

- (Pseudo) random numbers are key to many algorithms
 - a number of high-quality algorithms exist
- Typically generate number in the range $[0.0, 1.0)$
 - includes 0.0 but excludes 1.0
 - e.g. generate “random” integer i in range $0 \dots N-1$, return i / N
- Are often transformed to other distributions
 - analytically or using accept-reject stage
- Repeatability is a key requirement
 - necessary to test correctness of any computation
 - or interrupt and resume a long-running simulation